

# A study on quantifying effective training of DLDMD

---

A Thesis  
Presented to the  
Faculty of  
San Diego State University

---

In Partial Fulfillment  
of the Requirements for the Degree  
Master of Science in Applied Mathematics  
with a Concentration in  
Dynamical Systems

---

by  
Joseph Aaron Gene Diaz

July 12, 2022

# SAN DIEGO STATE UNIVERSITY

The Undersigned Faculty Committee Approves the  
Thesis of Joseph Aaron Gene Diaz:

A study on quantifying effective training of DLDMD

---

Christopher Curtis, Chair  
Department of Mathematics and Statistics

---

Jérôme Gilles  
Department of Mathematics and Statistics

---

Saeed Manshadi  
Department of Electrical Engineering

---

Approval Date

Copyright © 2022  
by  
Joseph Aaron Gene Diaz

## DEDICATION

Everyone who's helped me get this far: You know who you are.

## **ABSTRACT OF THE THESIS**

A study on quantifying effective training of DLDMD

by

Joseph Aaron Gene Diaz

Master of Science in Applied Mathematics with a Concentration in Dynamical Systems

San Diego State University, 2022

Write this once everything else is written.

# TABLE OF CONTENTS

	PAGE
ABSTRACT .....	v
LIST OF TABLES .....	vii
LIST OF FIGURES .....	viii
CHAPTER	
1 INTRODUCTION .....	1
1.1 What is E-DMD? .....	1
1.2 What are Neural Networks? .....	3
1.3 DLDMD and it's limitations .....	5
1.4 Statement of Dilemma .....	6
2 KULLBACK-LEIBLER DIVERGENCE .....	8
2.1 Basic Definitions .....	9
2.2 How to build distributions (Non-parametric statistics) .....	11
2.3 Implementation Details .....	11
3 RESULTS .....	13
3.1 Duffing .....	13
3.2 Van der Pol .....	13
3.3 Lorenz (Maybe) .....	14
4 DISCUSSION .....	17
BIBLIOGRAPHY .....	18

## LIST OF TABLES

	PAGE
3.1 describes the hyperparameters used for the Duffing Oscillator that were left unchanged for each run. ....	13

## LIST OF FIGURES

		PAGE
2.1	Entropy and density function change with respect to $\sigma^2$ .....	9
2.2	KL Divergence example with Normal distributions. As expected, the .....	10
3.1	Duffing average slopes and loss plots. ....	14
3.2	Duffing Average bandwidth. ....	15
3.3	Van der Pol average slopes and loss plots. ....	15
3.4	Van der Pol Average bandwidth. ....	16



# CHAPTER 1

## INTRODUCTION

In the study of dynamical systems a central problem is how to derive models from measured data to facilitate the prediction of future states. Many approaches and techniques exist in the literature, from deriving sets of governing equations via application of the simple physical principles to the statistically meaningful Principal Component Analysis and other modal decompositions. The main goal of many of these methods is to come up with a generalized framework so that the dynamics can be extracted from the data for the sake of control and prediction. While there are genuine advantages to the more concrete and scientifically-sound methods of deriving models from first principles, it is often very difficult if not outright untenable.

As such, data-driven methods have rapidly become very useful approaches for coming up with rough and ready models.

Say more once you've thought of it.

### 1.1 What is E-DMD?

As it says on the tin, we seek a predictive model for a time series  $\{\mathbf{y}_j\}_{j=1}^{N^T+1}$ , which are the measurements of a uniformly sampled unknown dynamical system of the form

$$\frac{d\mathbf{y}}{dt} = f(\mathbf{y}(t)), \quad \mathbf{y}(0) = \mathbf{x} \in \mathcal{M} \subseteq \mathbb{R}^{N_s}$$

As established in the introduction, there is an old and venerable literature dedicated to deriving system using classical methods. What this literature will admit is that such a process is not trivial and even difficult or impossible to do in practice; particularly for nonlinear phenomena worth investigating. We want something that is more easily generalized and algorithmic.

One method that can be leveraged for this is via the Koopman Operator  $\mathcal{K}$ . If we denote  $\boldsymbol{\varphi}(t; \mathbf{x}) = \mathbf{y}(t)$  to be the flow map affiliated with the initial condition  $\mathbf{x}$  and denote  $g : \mathcal{M} \rightarrow \mathbb{C}$  to be a square integrable observable of the system then, according to [7], there exists a linear representation of the flow map given by  $\mathcal{K}$ ;

$$\mathcal{K}^t g(\mathbf{x}) = g(\boldsymbol{\varphi}(t; \mathbf{x})),$$

which means that  $\mathcal{K}$  is a linear time-advancement operator of the dynamics. This might seem like it trivializes the problem, given that we now have a linear system, but it does

not. The Koopman operator is an *infinite* dimensional operator, so we've traded a potentially nonlinear problem for an infinite dimensional linear one.

In order to actually use this new representation, it suffices to find the eigenvalues  $\{\lambda_\ell\}$ , and their affiliated eigenfunctions  $\{\phi_\ell\}$ , of the Koopman Operator such that

$$\mathcal{K}^t \phi_\ell = \exp(t\lambda_\ell) \phi_\ell \implies g(\mathbf{x}) = \sum_{\ell \in \mathbb{N}} c_\ell \phi_\ell(\mathbf{x}),$$

where we can essentially construct a modal decomposition of  $g$ . From here advancing the dynamics to time  $t$  is equivalent to writing

$$\mathcal{K}^t g(\mathbf{x}) = \sum_{\ell \in \mathbb{N}} a_\ell \phi_\ell(\varphi(t, \mathbf{x}))$$

The most useful property of this framing is that we have a global linearization of the flow, with a major caveat: Generally, finding these eigenvalues and eigenfunctions is impossible. This led to the development of the already mentioned and much much lauded Dynamic Mode Decomposition (DMD) and its extensions, which seeks to numerically approximate a finite number of these modes. We'll focus on the particular extension that [1] implements: Extended DMD (E-DMD) [10]. With E-DMD, we suppose that

$$g(\mathbf{x}) = \sum_{\ell=1}^{N_O} a_\ell \psi_\ell(\mathbf{x})$$

which is to say that the observable  $g$  exists in a finite dimensional subspace of  $L^2(\mathcal{M})$ , applying the Koopman operator implies that

$$\begin{aligned} \mathcal{K}^{\delta t} g(\mathbf{x}) &= \sum_{\ell=1}^{N_O} a_\ell \phi_\ell(\varphi(\delta t, \mathbf{x})) \\ &= \sum_{\ell=1}^{N_O} \phi_\ell(\mathbf{x}) (\mathbf{K}_O^T \mathbf{a})_\ell + r(\mathbf{x}; \mathbf{K}_O) \end{aligned}$$

for discrete time step  $\delta t$ .  $\mathbf{K}_O$  is the  $N_O \times N_O$  matrix that minimizes

$$\mathbf{K}_O = \underset{K}{\operatorname{argmin}} ||\Psi_+ - K\Psi_-||_F^2$$

and  $r(\mathbf{x}, \mathbf{K}_O)$  is a residual that represents the total error due to DMD. If the ansatz that  $g$  lives in a finite space holds, then  $r$  is identically 0. We define  $\Psi_\pm$  to be

$$\Psi_- = (\Psi_1 \ \Psi_2 \ \cdots \ \Psi_{N_T}), \quad \Psi_+ = (\Psi_2 \ \Psi_3 \ \cdots \ \Psi_{N_T+1})$$

where  $\{\Psi_j\}$  is an observable of the time series of interest  $\{\mathbf{y}_j\}$ . What the expression for  $\mathbf{K}_O$  tells us is that, we are trying to find a one-step mapping from each data point to the next. Practically speaking,  $\mathbf{K}_O$  is found using an SVD, with which we can write

$$\Psi_- = U\Sigma W^\dagger \implies \mathbf{K}_O = \Psi_+ W \Sigma^{-P} U^\dagger$$

where  $-P$  denotes the Moore-Penrose pseudo-inverse and  $\dagger$  denotes the conjugate transpose. This gives us an expression for  $r$  in terms of the observables  $\Psi$ :

$$E_r(\mathbf{K}_O) = \|\Psi_+(I - WW^\dagger)\|_F$$

Finding the eigenvalues, eigenfunctions, and Koopman modes comes down to an eigen-decomposition, from which the dynamics can be approximated as

$$y(t; \mathbf{x}) \approx V \exp(t\Lambda) V^{-1} \Psi(\mathbf{x})$$

where  $\mathbf{K}_O = VTV^{-1}$ ,  $\Lambda_{\ell\ell} = \ln(T_{\ell\ell})/\delta t$  is a diagonal matrix and  $\Psi$  is the representation of the initial condition in terms of the observables.

**Expand on everything else above for a more complete picture of E-DMD.**

## 1.2 What are Neural Networks?

Before we move on, a brief digression on Neural Networks (NNs) is appropriate. Ever since scientists discovered the relatively simple interaction between neurons and axons in the human brain, they have been enamored with the ability to create a learning computer with the similar ability to become more adept at a task with training and practice. In the same way that art imitates life, the most straight-forward attempts have been to construct artificial neural networks that pantomime the biological ones that we carry around in our heads; and while these pale imitations have not been developed to rival the human brain, they have led to some major achievements in automation. In much the same way that a person is “trained” to perform a task by repeating it with feedback and adjusting their performance as they go, an artificial NN uses data, a loss function, and an optimization algorithm to change the state of the NN to one which can better accomplish the task. The loss function and optimizer are to the feedback and behavioral adjustment what the data is to the experience.

The oldest known example of NNs were the perceptrons of the 1950s, which could perform binary classifications of linearly separable data. Multilayer versions of this simple architecture allowed for classification of larger set of classes, but the researcher of the time arrived at the limits of their computing environments rather

quickly and a brief dark age in the study and application of neural networks ensued. Interest was renewed in the 1980s when many of the hardware limitations of previous decades were ameliorated and new algorithms for optimizing had been pioneered.

**Expand once you've got more to say.**

In modern mathematical study and application, the question of function representation is an ever-present one; specifically with respect to approximating quantities that can be given via series. While truncations and use of the likes of fourier series are useful and effective in some narrow applications, one means of representation that has shown much promise in recent years due to the greater availability of computing power is through the use of Neural Networks (NNs). In 1989, [6] showed that, with some appropriate parameters, a NN can be used to approximate any continuous function from one vector space to another to any arbitrary precision required.

From this we can consider ML as a tool of regression. Say we have a set of training data  $d_{tr} = \{(x_j, y_j)\}_{j=1}^{N_{tr}}$  and we believe that there is some function  $f(x)$  such that

$$y_j = f(x_j) + \varepsilon_j,$$

where  $\varepsilon_j$  is assumed to be a sample drawn from  $\mathcal{N}(0, \bar{\sigma}^2)$ . Generally  $f$  is not known; so we must, using the training data, build an estimator  $f_e(x; d_{tr})$  and measure the performance of the estimator by examining the quantity

$$\mathbb{E}_{d_{tr}} [(f(x) + \varepsilon - f_e(x; d_{tr}))^2] = \text{bias}_{d_{tr}}^2 + \mathbb{V}(f_e) + \bar{\sigma}^2$$

where

$$\text{bias}_{d_{tr}} = \mathbb{E}_{d_{tr}} [f_e] - f(x)$$

and

$$\mathbb{V}(f_e) = \mathbb{E}_{d_{tr}} [(\mathbb{E}_{d_{tr}} [f_e] - f_e(x; d_{tr}))^2]$$

With this framing, ML methods can be cast as attempts to find an estimator relative to a given training that balances the competing dilemmas of being too biased or allowing too much variance. From the point of view of non-parametric statistics, we would need to utilize other techniques[5] to determine the bias and variance of any given estimator. For our purposes, we will instead consider how metrics from information theory can be used to assess the change in information as data is fed through a NN and a model is trained.

### 1.3 DLDM and its limitations

**Incomplete picture, add more detail later. Consider re-skimming Jay's paper.**

The key innovation of [1] is to use a neural network to come up with the collection of observables on  $\{\mathbf{y}_j\}$  that allow for the best prediction of future system states. This is implemented by defining an encoder  $\mathcal{E} : N_S \rightarrow N_O$  and decoder  $\mathcal{D} : N_O \rightarrow N_S$  composed of Dense layers such that

$$(\mathcal{D} \circ \mathcal{E})(\mathbf{x}) = \mathbf{x}$$

We choose  $N_O \geq N_S$  and an appropriate loss function so that  $\mathcal{E}$  and  $\mathcal{D}$  give a richer space of observables, called the latent space, for EDMD to use when advancing the dynamics. The implementation of NNs for this purpose requires a method of tuning to allow  $\mathcal{E}$  and  $\mathcal{D}$  to learn the best representations possible. As such, a loss function that correctly identifies and prioritizes the desired properties is a necessary condition for the DLDM to function as needed. A natural choice considering these constraints is given by

$$\mathcal{L} = \alpha_1 \mathcal{L}_{\text{recon}} + \alpha_2 \mathcal{L}_{\text{dmd}} + \alpha_3 \mathcal{L}_{\text{pred}} + \alpha_4 \|\mathbf{W}_g\|_2^2 \quad (1.1)$$

where

$$\mathcal{L}_{\text{recon}} = \frac{1}{N_T + 1} \sum_{j=1}^{N_T+1} \|\mathbf{y}_j - (\mathcal{D} \circ \mathcal{E})(\mathbf{y}_j)\|_2, \quad (1.2)$$

$$\mathcal{L}_{\text{dmd}} = E_r(\mathbf{K}_O), \quad (1.3)$$

$$\mathcal{L}_{\text{pred}} = \frac{1}{N_T} \sum_{j=1}^{N_T} \|\mathbf{y}_{j+1} - \mathcal{D}(VT^j V^{-1} \mathcal{E}(x))\|_2, \quad (1.4)$$

Each component guides the machine to a particular outcome:

1.  $\mathcal{L}_{\text{recon}}$  is the Mean Squared Error (MSE) of each time step with respect to the reconstruction from the composition of  $\mathcal{E}$  and  $\mathcal{D}$ . This component ensures that, under training, the network effectively acts as a near identity transformation for data that is fed into it. This quality allows the DMD advanced trajectories to be recovered from the higher dimensional latent space back to the original dimension of the data.
2.  $\mathcal{L}_{\text{DMD}}$  is the error associated with DMD. Consequently, this component is the one that is most responsible for finding the optimal set of observables for DMD.  $\mathcal{E}$  immerses the data into a higher dimensional latent space, in effect acting as our set of observable; minimizing this gives us greater flexibility in the latent space.
3.  $\mathcal{L}_{\text{pred}}$  is the (MSE) for each forward time prediction due to DMD and immersion/submersion due to  $\mathcal{E}/\mathcal{D}$ . In addition to balancing the last conditions,

this condition ensures that the DMD step in the latent dimension is consistent with the next time-step from the time series after encoding, advancing, and decoding.

4.  $\mathbf{W}_g$  is the vectorized quantity that represents all weights in both  $\mathcal{E}$  and  $\mathcal{D}$ . This is really only a regularization condition to keep the coefficients of the weight matrices from blowing up in value as the model trains, which can be a concern for ML models.
5.  $\alpha_1$ ,  $\alpha_2$ ,  $\alpha_3$ , and  $\alpha_4$  are 4 positive constants that allow us to assign a weighting to each component of the loss. This allows the loss function to be dynamically weighted to prioritize some conditions over others. In [1],  $\alpha_1 = \alpha_2 = \alpha_3 = 1$  and  $\alpha_4 < 10^{-10}$ .

The marriage of E-DMD with Dense NNs and some meaningful choices of loss function as an innovation was made by [1, 3, 8], and has proved to be quite a robust method for learning dynamics from data. In researching this thesis, many attempts were made to test the limits of DLDMD. [1] was primarily interested in recreating the phase space behavior of a handful of well known nonlinear oscillators such as the nonlinear harmonic oscillator, the Duffing oscillator, and the Van der Pol oscillator. Additionally, chaotic systems like Lorenz 63 and the Rössler system were examined but chaos proved to be quite a challenge for the algorithm to overcome. A study into the addition of gaussian noise led to the discovery of the machine being quite robust to data due to the characteristics of it's loss function and additional noise mitigation was accomplished by implementing convolutional layers instead of traditional dense layers.

## 1.4 Statement of Dilemma

Having established DLDMD and having discussed testing it's limitations via training new models and *visually* qualifying each run as a success or failure, a question is begged: Can we, in an empirical and quantitative fashion, determine whether good training is taking place or otherwise deduce whether our model is well-posed given the hyperparameters? A machine learning model that must learn how to accomplish even a simple task can take quite a long time and plenty of computing resources to complete it's training and only then can the results be checked to determine whether productive training was taking place as opposed to the optimizer becoming stuck in a local minimum or there being no global minimum at all.

How can this be quantified? A useful starting point would be to examine the weights of the matrices that make up the layers of a model. For Dense layers, passing a vector of data  $\mathbf{x} \in \mathbb{R}^d$  through a layer  $L$  can be written as

$$L(\mathbf{x}; A, \sigma, b) = \sigma(A\mathbf{x} + b)$$

for some matrix  $A \in \mathbb{R}^{D \times d}$ , vector  $b \in \mathbb{R}^D$  and (typically nonlinear) activation function  $\sigma : \mathbb{R}^D \rightarrow \mathbb{R}^D$ . For a given layer,  $\sigma$  is a fixed hyperparameter of the machine, but both  $A$  and  $b$  have entries that are tuned by the optimizer as training takes place. Training takes place on a per epoch basis for most models, so considering the configuration of these matrices epoch to epoch can help us determine whether they are converging to a set of elements that accomplish the desired task or not. The next state for each layer is dependent on the previous, so a reasonable framing would be as one of a discrete dynamical system of the form

$$Q_{n+1} = P(Q_n) \quad (\text{Come up with better symbols later})$$

where  $P$  is the training procedure and  $Q_n$  represents the configuration of the network's weights at each epoch  $n$ . Typically, the initial state  $Q_0$  is “random” in the sense that the weights are selected from a probability distribution. In the interest of examining convergences, one might consider the following limit

$$\lim ||Q_{n+1} - Q_n||_2$$

for some  $L^2$  norm on the space that  $Q_n$  inhabits. While this “one-step” Cauchy convergence will tell us whether the machine is approaching a particular configuration point-wise, this might not actually tell us much of anything else. The changes between epochs is, in some sense, stochastic due to not knowing how the optimizer will update  $Q$  between epochs. As such, we propose a more statistical approach wherein we examine how the information content is changing from epoch to epoch. By the noted stochastic nature of the evolution, we consider the entries of  $A$  and  $b$  for each layer to be drawn from a continuous probability distribution  $X$  and examine how it changes epoch to epoch.

## CHAPTER 2

### KULLBACK-LEIBLER DIVERGENCE

In the field of information theory, it is often useful to frame information as being represented by a random variable and the values it can take on. For example, consider a normally distributed random variable  $X \sim \mathcal{N}(\mu, \sigma^2)$  where  $\mu$  is the mean or expected value  $\mathbb{E}[X]$  of  $X$  and  $\sigma^2$  is the variance  $\mathbb{V}[X]$  of  $X$ . This is one of the most well understood probability distributions due to countless years of study into its properties, as such we know that if we were to draw a sample for  $X$  that we're much more likely to get something close to  $\mu$  than we are to get something far away from  $\mu$ . For example; if  $\mu = 0$  and  $\sigma^2 = 1$ , then upon drawing a sample the value 1 is more likely to appear than -37. The level of certainty of the outcome of the draw begs the question of how that certainty can be quantified and conversely, we could instead consider the uncertainty of any given probability distribution. A measure of the uncertainty was proposed in 1948 and this measure is known as *informational entropy* [9] or, simply, entropy. For a continuous distribution with density function  $f(x)$ , the entropy is given by

$$h[f] = \mathbb{E}[-\log(f(x))] = - \int_{\mathcal{X}} f(x) \log(f(x)) \, dx$$

For the normal distributed  $X$  defined above, we have that probability density function

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2\right) \quad (2.1)$$

and the entropy integral for this evaluates to


$$h[f] = \frac{1}{2} (\log(2\pi\sigma^2) + 1)$$

which only depends on the variance. This expression tells us that the entropy of a normally distributed random variable increases as the variance does, graphically this is associated with the flattening of the bell curve which can be interpreted as the values around the mean being less likely to appear relative to values far from the mean. As such, the result of any given sample is less certain.

#### **Bias vs Variance.**

When training an ML model, the expectation is that the machine will converge to some configuration that allows it to accomplish the task it was designed for. For the sake of





**Figure 2.1.** Entropy and density function change with respect to  $\sigma^2$ .

## 2.1 Basic Definitions

The KL Divergence (KLD) is a statistical distance between a pair of probability distributions which measures how different a distribution  $P$  is from a reference distribution  $Q$ . “A simple interpretation of the divergence of  $P$  from  $Q$  is the expected excess surprise from using  $Q$  as a model when the actual distribution is  $P$ .” If  $P$  and  $Q$  are discrete distributions defined on the probability space  $\mathcal{X}$ , then the KLD of  $P$  with respect to  $Q$  is given by

$$D_{KL}(P \parallel Q) = \sum_{x \in \mathcal{X}} P(x) \log(P(x)/Q(x))$$

In other words, it is the expectation of the logarithmic difference between the probabilities  $P$  and  $Q$ , where the expectation is taken using the probabilities  $P$ . As a “distance”, the KLD is non-negative and 0 when  $P = Q$  almost everywhere. Unlike more standard metrics, it is not symmetric and does not satisfy the triangle equality. In order to deal with something that is, at least, symmetric we’ll consider using the symmetric KL Divergence which is defined as

$$D_{SKL}(P, Q) = \frac{1}{2} (D_{KL}(P \parallel Q) + D_{KL}(Q \parallel P))$$

Which is the average of the KLD for  $P$  with respect to  $Q$  and for  $Q$  with respect to  $P$ . In our cases, we’ll be dealing with a continuous random variable with probability density functions  $p$  and  $q$ , for which the corresponding KLD formula is

$$D_{KL}(P \parallel Q) = \int_{\mathcal{X}} p(x) \log(p(x)/q(x)) \, dx$$

As an analogue of the discrete case, the same simple interpretations about the formula apply. The reason we can consider this a “distance” is that the KLD is non-negative,


$D_{KL}(P \parallel Q) \geq 0$ , which is known as Gibbs inequality and only equals 0 when  $P = Q$  almost everywhere. As such, the distance between  $P$  and  $Q$  is 0 when they are the “same” distribution and some positive number if they differ. For example, consider the random variables  $X_1 \sim \mathcal{N}(\mu_1, \sigma_1^2)$  and  $X_2 \sim \mathcal{N}(\mu_2, \sigma_2^2)$  with corresponding density functions given by equation 2.1; the divergence between the distributions is

$$D_{KL}(p \parallel q) = \frac{1}{2} \left( \ln \left( \frac{\sigma_2^2}{\sigma_1^2} \right) + \frac{\sigma_1^2}{\sigma_2^2} + \frac{(\mu_1 - \mu_2)^2}{\sigma_2^2} - 1 \right), \quad (2.2)$$

Letting  $\mu_1 = \mu_2$ ,  $\sigma_1^2 = 1$ , and allowing  $\sigma_2^2$  to vary, we can write

$$D_{KL}(p \parallel q) = \frac{1}{2} \left( 2 \ln \sigma_2 + \frac{1}{\sigma_2^2} - 1 \right), \quad (2.3)$$

Figure 2.2 illustrates



**Figure 2.2.** KL Divergence example with Normal distributions. As expected, the

De-trending. KL div of differences.

## 2.2 How to build distributions (Non-parametric statistics)

The question remains about how we obtain the probability distribution of the random variable that represents the evolution of each layer of an ML model from measured data. Many techniques exist, but we'll be using Kernel Density Estimation (KDE). The most basic idea behind KDE is using a different bin centering method and a smoothing factor, called a kernel, to approximate a probability density  $f$  from measured data in the form of a histogram. This is accomplished with the *kernel density estimator* for  $f$  given by:

$$\hat{f}(x; K, h) = \frac{1}{nh} \sum_{i=1}^n K\left(\frac{x - x_i}{h}\right)$$

where  $K$  is the chosen kernel,  $h$  is what's called the bandwidth parameter, and  $\{x_i\}_{i=1}^n$  is the data that generates our histogram. Per [4], the choice of kernel does not provide much statistical significance; the choice of the bandwidth parameter  $h$  is very crucial for finding a density estimate that approximates the underlying density function appropriately. Choosing an optimal  $h$  depends on the given data and the common method, Silverman's rule of thumb, works on the assumption that the density function of the data is unimodal and close to normal. For general data, on which no assumptions of normality are made, a more general method can be found using the Improved Sheather-Jones algorithm[2].

## 2.3 Implementation Details

Ok, so what the hell did we actually do?

1. Ran a series of DLDMD models on two very well understood dynamical systems:
  - Van der Pol oscillator and Multiple nonlinear centers Duffing oscillator
  - Used latent/lifting dimensions 2 - 15
  - Trained for 1000 epochs
  - Saved configuration of DLDMD network weights at each epoch
  - Used different data for each training, but each data set was drawn from the same system's dynamics. Working on reproducing with same data/determining if that's necessary.
2. Detrended weight data via differencing and used Kernel Density Estimation to create empirical PDFs from inter-epoch weights.
  - For sake of ensuring statistical significance, kernel and bias weights are combined.

- Used Epanechnikov kernel (optimal in a mean square error sense) which is  $K(x) = 3(1 - x^2)/4$  for  $|x| < 1$  else  $K(x) = 0$  and has finite support [4].
  - Kernel choice not stastically significant [4].
  - Bandwidth computed using Improved Sheather-Jones method [2] and saved for each epoch and layer.
3. Used symmetric KL Divergence to find “distance” between empirical PDFs from data.
    - Both PDFs are assumed to have same support so that integration is over same domain.
    - Using Simpson’s Rule with 10,000 points for numerical integration, considering approximating fourth derivative for error bound on integration.
  4. Using Divergence data for each layer and inter-epoch, found linear fit with 95% confidence and prediction interval.
    - First 20% of divergence data is ignored to avoid transient behavior.
    - Slope of each layer and model saved, computed average and variance of slopes among all layers.
  5. Plotted average slopes with variances of fit for each latent dimension.
    - Plots of the bandwidth per layer, inter-epoch, and latent dimension were made along with average bandwidths per latent dimension and inter-epoch.
    - Plots of all loss curves per model and latent dimension were also made.

## CHAPTER 3

### RESULTS

From experimentation by [1], we know that some sets of hyperparameters yield good reconstruction of the structure of the phase space of each the dynamical systems in question and that some do not.

#### 3.1 Duffing

**Table 3.1.** describes the hyperparameters used for the Duffing Oscillator that were left unchanged for each run.

Hyperparameter	Value
Initial Conditions	5000
Training proportion	66.66%
Validation proportion	19.98%
Testing proportion	13.36%
$t_0$	0
$\Delta t$	0.05
$t_f$	20
Epochs	1000
Encoder layers	3
Encoder hidden activation	elu
Encoder output activation	identity
Decoder layers	3
Decoder hidden activation	elu
Decoder output activation	identity
Neurons/layer	128
Weight initializer	$N(0, .1, 2)$
Bias Initializer	0
Learning Rate	$10^{-3}$
$a_1, a_2, a_3$	1
$a_4$	$10^{-14}$

#### 3.2 Van der Pol




Figure 3.1. Duffing average slopes and loss plots.

### 3.3 Lorenz (Maybe)




Figure 3.2. Duffing Average bandwidth.

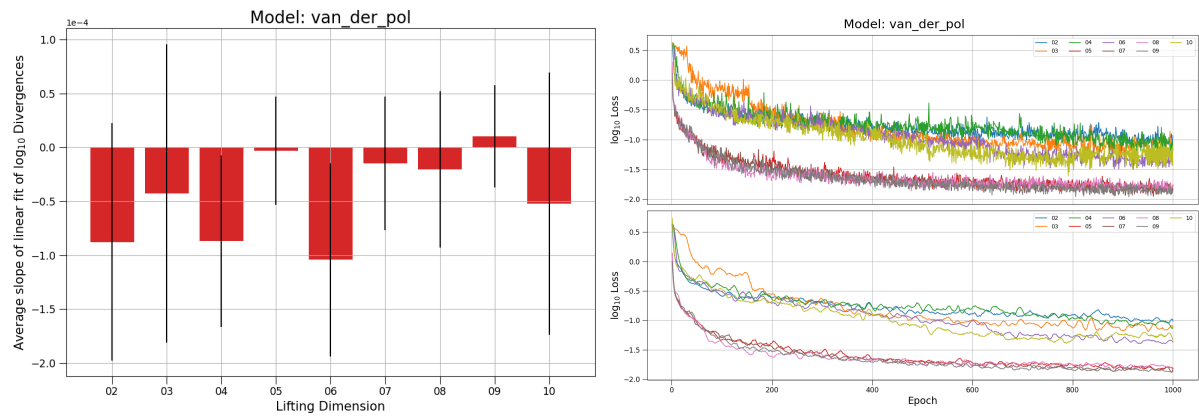


Figure 3.3. Van der Pol average slopes and loss plots.




Figure 3.4. Van der Pol Average bandwidth.



## CHAPTER 4

### DISCUSSION

## BIBLIOGRAPHY

- [1] D. ALFORD-LAGO, C. W. CURTIS, A. T. IHLER, AND O. ISSAN, *Deep Learning Enhanced Dynamic Mode Decomposition*, (2022).
- [2] Z. I. BOTEV, J. F. GROTHOWSKI, AND D. P. KROESE, *Kernel density estimation via diffusion*, The Annals of Statistics, 38 (2010), pp. 2916 – 2957.
- [3] S. BRUNTON, B. R. NOACK, AND P. KOUMOUTSAKOS, *Machine learning for fluid mechanics*, Ann. Rev. Fluid Mech., 52 (2020), pp. 477–508.
- [4] V. A. EPANECHNIKOV, *Non-parametric estimation of a multivariate probability density*, Theory of Probability & Its Applications, 14 (1969), pp. 153–158.
- [5] M. HOLLANDER, D. WOLFE, AND E. CHICKEN, *Nonparametric Statistical Methods*, Wiley Series in Probability and Statistics, Wiley, 2013.
- [6] K. HORNIK, M. STINCHCOMBE, AND H. WHITE, *Multilayer Feedforward Networks are Universal Approximators*, Neural Networks, (1989).
- [7] B. KOOPMAN, *Hamiltonian systems and transformations in Hilbert space*, Proc. Nat. Acad. Sci., 17 (1931), pp. 315–318.
- [8] B. LUSCH, J. N. KUTZ, AND S. L. BRUNTON, *Deep learning for universal linear embeddings of nonlinear dynamics*, Nature Comm., 9 (2018), p. 4950.
- [9] C. E. SHANNON, *A mathematical theory of communication*, The Bell System Technical Journal, 27 (1948), pp. 379–423.
- [10] M. WILLIAMS, I. G. KEVREKIDIS, AND C. W. ROWLEY, *A data-driven approximation of the Koopman operator: extending dynamic mode decomposition*, J. Nonlin. Sci., 25 (2015), pp. 1307–1346.